

# Michael Sparre

[micsparre@gmail.com](mailto:micsparre@gmail.com)

[michaelsparre.com](https://michaelsparre.com)

[github.com/micsparre](https://github.com/micsparre)

[linkedin.com/in/micsparre](https://linkedin.com/in/micsparre)

## Work Experience

---

**Founding Engineer**, Grepr — San Francisco, CA

Jan 2025 — Present

- Owned the design, implementation, and rollout of a Kubernetes operator replacing a 500 ms server-side polling control plane; migrated every production Flink pipeline, rolled the operator across five environments, supported 1,300+ reconfigurations in 60 days without a failed reconfiguration, and retired the legacy reconciler.
- Refactored a multi-tenant impact-estimation platform processing 90K+ imported monitors, dashboards, and log metrics; deduplicated 53K estimation records into 5.1K Flink batch computations, reducing required batch work by 10x and vendor API fan-out by 2x.
- Led two customer POCs through telemetry-sink migration, pipeline deployment, and reduction tuning; both became paying production customers and now sustain 90–93% log-volume reduction across 14B+ monthly events.
- Designed an extensible attribute-aggregation framework for Grepr’s Flink log reducer, shipping backend configuration and UI while increasing global-reducer throughput by 40%.

**Software Engineer**, Gilead Sciences — Bay Area, CA

May 2021 — Dec 2024

- Built the department’s first observability platform with Grafana and Prometheus, enabling proactive monitoring across AWS data-lake pipelines and PostgreSQL data products.
- Automated 10+ ingestion pipelines that consolidated scientific and operational data into trusted master-data products used by downstream analytics and applications.

## Projects

---

**workout**

[workout.michaelsparre.com](https://workout.michaelsparre.com)

- Built and deployed a multi-user Hevy training platform with webhook-driven sync, progress analytics, and a scoped MCP server for AI-assisted workout analysis and safe routine updates.

**observe**

[observe.michaelsparre.com](https://observe.michaelsparre.com)

- Built and deployed a self-hosted personal-finance platform with Plaid, SnapTrade, React, FastAPI, and PostgreSQL, plus portfolio and budget analytics, MCP access, and production observability.

## Skills

---

**Languages:** Java, Python, SQL, TypeScript, Bash

**Infrastructure:** Kubernetes, AWS, Docker, Terraform, Helm, CI/CD, Prometheus, Grafana

**Systems & Tools:** Apache Flink, Kafka, PostgreSQL, Iceberg, React, LLM tool calling, MCP, RAG

## Education

---

**University of California, Berkeley**

May 2022

B.S. in Electrical Engineering and Computer Sciences